CTS SRE – 17

Incident Simulation and Root Cause Analysis (RCA)

EMPOWERING High Performance Technology Teams

Agenda for the Day

- Introduction to Incident Simulation & RCA
- Simulating Production Incidents
- Incident Lifecycle Management
- Advanced Root Cause Analysis (RCA) Techniques
- Predictive Maintenance & AI-Driven Failure Prediction
- GenAI-Powered RCA: Case Study on Payment Gateway Latency
- Postmortem Report Best Practices & Stakeholder Communication
- Summary & Key Takeaways

Introduction to Incident Simulation & RCA

What is Incident Simulation?

- Controlled execution of failure scenarios to test system resilience.
- Helps improve team coordination, debugging skills, and response time.

Why is Root Cause Analysis (RCA) Important?

- Prevents recurring failures by identifying the actual cause of an incident.
- Improves system reliability and reduces downtime.

Key Objectives:

- Enhance real-time incident handling capabilities.
- Develop **systematic troubleshooting methodologies**.
- Leverage **AI & machine learning for predictive failure analysis**.

Simulating Production Incidents

Types of Incident Simulations:

- **Application-Level Failures** (API errors, service crashes, memory leaks)
- **Infrastructure Failures** (Server crashes, disk failures, network partitions)
- **Database Issues** (Deadlocks, slow queries, replication lag)
- **Third-Party API Failures** (Payment gateway issues, slow external services)

Simulation Tools & Techniques:

- **Chaos Engineering:** AWS Fault Injection Simulator (FIS), Gremlin, LitmusChaos
- **Traffic Surge Testing:** Load testing with JMeter, Locust
- **Network Fault Injection:** TC (Linux Traffic Control), Toxiproxy

🛠 **Example:** Simulating a regional server failure and measuring auto-recovery time.

Incident Lifecycle Management

Phases of Incident Lifecycle:

Detection: Identifying system anomalies and failures.

Diagnosis: Gathering logs, metrics, and traces for troubleshooting.

Containment: Mitigating immediate impact and preventing escalation.

Resolution: Applying fixes, rolling back deployments, or scaling resources.

Postmortem & Continuous Improvement: Learning from failures and improving resilience.

Key Monitoring & Alerting Tools:

AWS CloudWatch, Prometheus, ELK Stack, Datadog

PagerDuty, Opsgenie, Slack for real-time alerts

🛠 **Example:** Handling a critical **database connection failure** during peak traffic

Advanced Root Cause Analysis (RCA) Techniques

Effective RCA Frameworks:

- **5 Whys Methodology** – Ask "Why?" multiple times to trace failure origins.
- **Ishikawa (Fishbone) Diagram** – Categorize failure causes by system components.
- **Fault Tree Analysis (FTA)** – Logical breakdown of root causes.

Analyzing Multi-Layered System Failures:

- **Application Failures** – Debugging logs, tracing API calls, identifying bottlenecks.
- **Infrastructure Issues** – Analyzing CPU, memory, network bottlenecks.
- **Third-Party Dependencies** – Latency issues in API calls or payment gateways.

Best Practices for Stakeholder Communication:

- Maintain **transparency** with business teams and leadership.
- Use **data-driven insights** for explaining failures and corrective actions.
- 🛠 **Example:** Conducting an RCA for a **slow-loading e-commerce checkout page**.

Predictive Maintenance & AI-Driven Failure Prediction

What is Predictive Maintenance?

Analyzing historical system data to predict and prevent failures.

AI-Powered Anomaly Detection:

Identifying patterns in system logs, traces, and performance metrics.

Using ML algorithms to detect abnormal system behaviors.

Key Predictive Analytics Tools:

AWS DevOps Guru, Datadog AIOps, Splunk Machine Learning Toolkit

🛠 **Example:** Predicting **disk failure in a high-throughput database system** using anomaly detection.

GenAI-Powered RCA – Case Study on Payment Gateway Latency

Scenario: A global e-commerce company experiences a **high failure rate in payment transactions**, leading to revenue loss.

Challenges Identified:

Payment gateway response times fluctuate, causing delayed transactions.

API retries cause cascading failures, leading to **increased load on backend services**.

GenAI Solution:

Analyze **real-time transaction logs and response times**.

Generate **automated RCA reports** identifying performance anomalies.

Propose **dynamic API retry strategies** to reduce failure rates.

🛠 **Outcome:**

Reduced **payment failures by 40%** through intelligent traffic rerouting.

Optimized **retry logic to minimize API overload issues.**

Postmortem Report Best Practices & Stakeholder Communication

Importance of Postmortems:

Ensures **continuous learning** from incidents.

Helps **identify process gaps** and improve future incident response.

Key Components of a Postmortem Report:

Summary: High-level impact description.

Timeline: Chronology of events leading to failure.

Root Cause: Breakdown of technical failure points.

Lessons Learned: Preventative actions for future incidents.

Avoiding a Blame Culture:

Focus on **fixing systems, not blaming people**.

Encourage **collaborative problem-solving**.

🛠 **Example:** Postmortem of a **multi-region AWS S3 outage affecting customer file storage.**

Breakout

Summary & Key Takeaways

✔ Incident simulation enhances real-world troubleshooting skills. ✔ RCA helps prevent recurring failures and optimizes system reliability. ✔ Predictive maintenance reduces downtime and enhances proactive response. ✔ GenAI accelerates RCA by automating anomaly detection and failure prediction. ✔ A well-documented postmortem improves operational resilience.

Introduction to case studies

Expectation from case study

🚀 **Participants are expected to work on the deliverables:**

Define **SLI/SLO metrics** and create monitoring dashboards.

Automate cloud infrastructure provisioning and deployments.

Run **chaos experiments** and document resilience gaps.

Conduct **root cause analysis (RCA)** and propose improvements.

Thank You!

StackRoute

(An NIIT Venture)